

In [26]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

In [27]:

```
df=pd.read_csv(r"C:\Users\lenovo\Downloads\fiat500_VehicleSelection_Dataset.csv1.csv")
df
```

Out[27]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	
0	1	lounge	51	882	25000	1	44.907242	8.611
1	2	pop	51	1186	32500	1	45.666359	12.241
2	3	sport	74	4658	142228	1	45.503300	11.417
3	4	lounge	51	2739	160000	1	40.633171	17.634
4	5	pop	73	3074	106880	1	41.903221	12.495
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704
1534	1535	lounge	74	3835	112000	1	45.845692	8.666
1535	1536	pop	51	2223	60457	1	45.481541	9.413
1536	1537	lounge	51	2557	80750	1	45.000702	7.682
1537	1538	pop	51	1766	54276	1	40.323410	17.568

1538 rows × 9 columns

In [28]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null   int64
1   model          1538 non-null   object
2   engine_power    1538 non-null   int64
3   age_in_days     1538 non-null   int64
4   km             1538 non-null   int64
5   previous_owners 1538 non-null   int64
6   lat            1538 non-null   float64
7   lon            1538 non-null   float64
8   price          1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [29]:

```
df.describe()
```

Out[29]:

	ID	engine_power	age_in_days	km	previous_owners	li
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.54136
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.13351
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.85583
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.80295
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.39405
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.46796
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.79561

In [30]:

```
df.columns
```

Out[30]:

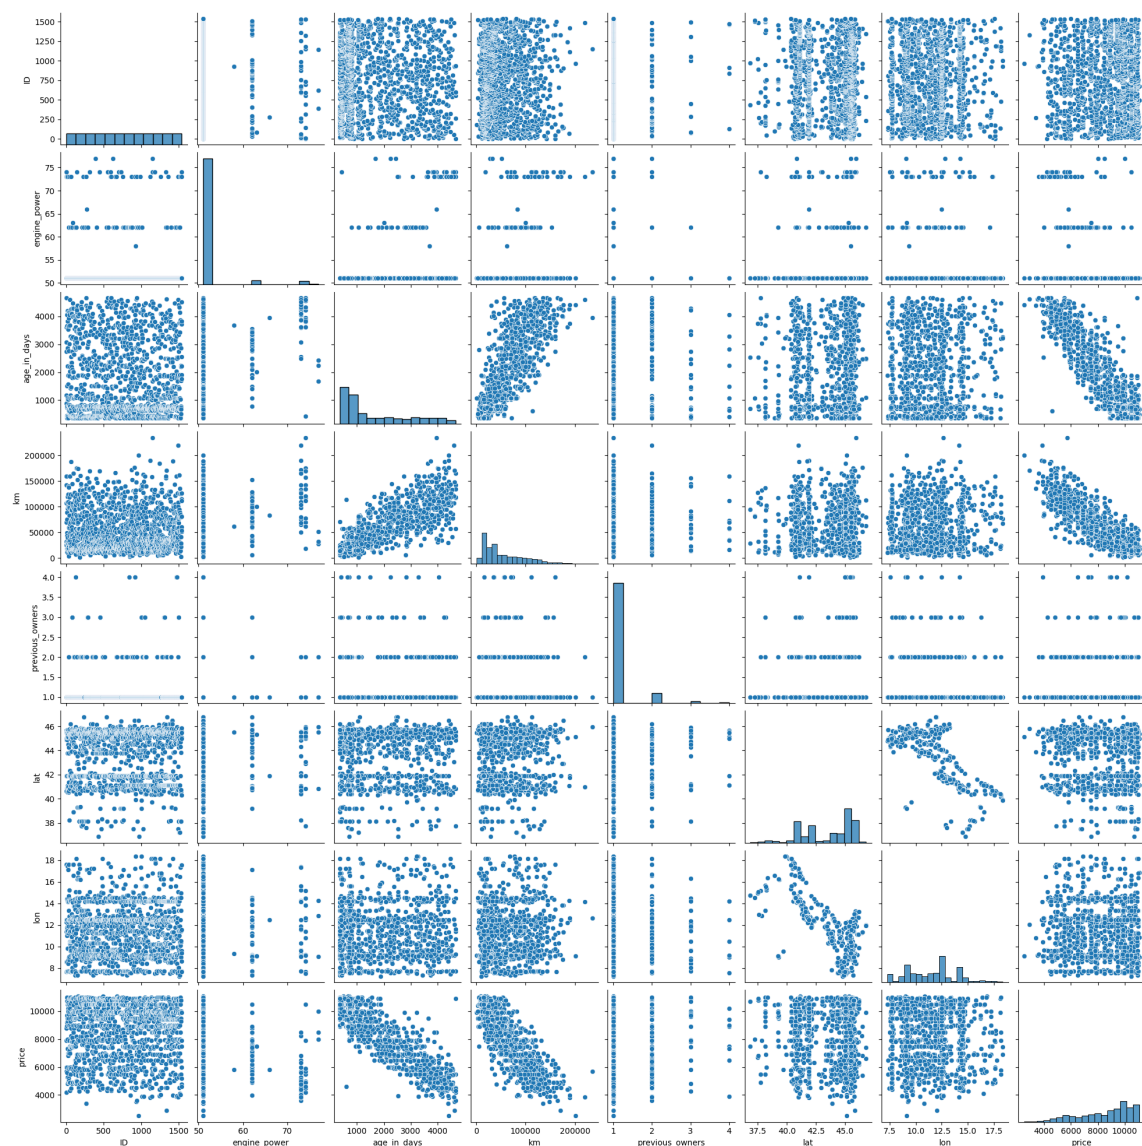
```
Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
      'lat', 'lon', 'price'],  
      dtype='object')
```

In [31]:

```
sns.pairplot(df)
```

Out[31]:

<seaborn.axisgrid.PairGrid at 0x19dee732950>




In [32]:

```
sns.displot(['Price'])
```

Out[32]:

<seaborn.axisgrid.FacetGrid at 0x19df1bab050>




In [33]:

```
sns.displot(df['price'])
```

Out[33]:

<seaborn.axisgrid.FacetGrid at 0x19df1eac1d0>




In [34]:

```
sns.displot(df['engine_power'])
```

Out[34]:

<seaborn.axisgrid.FacetGrid at 0x19ded8e5390>




In [35]:

```
fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
          'lat', 'lon', 'price']]  
sns.heatmap(fiatdf.corr())
```

Out[35]:

<Axes: >




In [36]:

```
fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',
           'lat', 'lon']]
sns.heatmap(fiatdf.corr())#without price
```

Out[36]:

<Axes: >



In [37]:

```
X=fiatdf[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',
           'lat', 'lon']]
y=df['price']
```

In [38]:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=101)
from sklearn.linear_model import LinearRegression
regr=LinearRegression()
regr.fit(X_train,y_train)
print(regr.intercept_)
```

8971.195683500027

In [39]:

```
coeff_df=pd.DataFrame(regr.coef_,X.columns,columns=['coefficient'])  
coeff_df
```

Out[39]:


	coefficient
ID	-0.046704
engine_power	11.646408
age_in_days	-0.898018
km	-0.017232
previous_owners	26.400886
lat	32.189709
lon	0.161073

In [40]:

```
predictions=regr.predict(X_test)  
plt.scatter(y_test,predictions)
```

Out[40]:

<matplotlib.collections.PathCollection at 0x19df3f98350>




In [41]:

```
sns.displot((y_test,predictions),bins=50)
```


Out[41]:

<seaborn.axisgrid.FacetGrid at 0x19dedc833d0>



In [42]:

```
sns.displot((y_test,predictions),bins=50);
```



In [43]:

```
from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 593.0876179519989

MSE: 551442.6799691883

MAE: 742.5918663500081

In [44]:

```
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.8597136704308846

In [45]:

```
df.fillna(method='ffill',inplace=True)
```

In [46]:

```
x=np.array(df['age_in_days']).reshape(-1,1)
y=np.array(df['km']).reshape(-1,1)
df.dropna(inplace=True)
```

In [47]:


```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

Out[47]:

```
▼ LinearRegression
LinearRegression()
```

In [48]:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



In [51]:

```
#elasticnet
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
[25.89689696]
[10640.73996329]
Mean Squared Error on test set 2690134778.1505136
```

In []: